MNIST Classification using PCA, FDA, and Discriminant Analysis

Rishabh Gupta
2024461

March 10, 2026

1 Introduction

In this assignment, we build a robust classification pipeline for handwritten digits (0, 1, and 2) from the MNIST dataset. Building upon Maximum Likelihood Estimation (MLE), this pipeline introduces dimensionality reduction techniques to handle the high-dimensional feature space (784 pixels) of the images. We implement and compare Principal Component Analysis (PCA) and Fisher's Discriminant Analysis (FDA) to project the data into lower dimensions before applying Linear Discriminant Analysis (LDA) and Quadratic Discriminant Analysis (QDA) for final classification.

2 Theoretical Background

2.1 Principal Component Analysis (PCA)

PCA is an unsupervised dimensionality reduction technique that seeks to maximize the variance of the projected data. Given a mean-centered data matrix X_c, we compute the covariance matrix $S = \frac{1}{N-1} X_c X_c^T$. By extracting the eigenvectors corresponding to the largest eigenvalues of S, we form a projection matrix U_p that retains a target percentage of the original variance.

2.2 Fisher's Discriminant Analysis (FDA)

Unlike PCA, FDA is a supervised technique that explicitly aims to maximize class separability. It achieves this by finding a projection matrix W that maximizes the ratio of the between-class scatter (S_B) to the within-class scatter (S_W):

$$W = \arg\max_W \frac{\det(W^T S_B W)}{\det(W^T S_W W)} \tag{1}$$

This translates to solving the generalized eigenvalue problem $S_W^{-1} S_B w = \lambda w$.

3 Implementation Details

3.1 Dimensionality Reduction: PCA

The PCA pipeline centers the data, calculates the covariance matrix, and dynamically selects the number of principal components needed to reach a specific variance threshold (e.g., 75% or 90%).

```
def get_pca(dataset, target_var):
    def get_pca(dataset, target_var):
    images_c = dataset

    N_c = len(images_c) - 1
```

```
    mu_list = compute_mean(images_c)

    images_c = np.array(images_c)
    mu = np.array(mu_list)

    images_centered = images_c - mu
    sigma = (images_centered.T @ images_centered) / N_c

    eigenvalues, eigenvectors = np.linalg.eigh(sigma)

    eigenvalues = np.real(eigenvalues)
    eigenvectors = np.real(eigenvectors)

    sorted_indices = np.argsort(eigenvalues)[::-1]
    eigenvalues = eigenvalues[sorted_indices]
    eigenvectors = eigenvectors[:, sorted_indices]

    total_var   = np.sum(eigenvalues)
    current_var = 0
    idx = -1

    for i in range(len(eigenvalues)):
        current_var += eigenvalues[i]
        if (current_var >= target_var*total_var):
            idx = i
            break

    return eigenvectors[:, :idx+1]
```

3.2 Class Projection: FDA

For FDA, we compute the global mean and class-specific means to build the S_B and S_W scatter matrices. A small regularization noise (10^{-5}) is added to the diagonal of S_W to ensure invertibility before solving for the eigenvectors.

```
def get_fda(images, label):
    classes = [0, 1, 2]

    global_mean = np.array(compute_mean(images))
    features = train_x.shape[1]
    S_w = np.zeros((features, features))
    S_b = np.zeros((features, features))

    for c in classes:
        images_c = []
        for i in range(len(label)):
            if(label[i] == c):
                images_c.append(images[i])
        N_c = len(images_c)

        mu = np.array(compute_mean(images_c))
        images_c = np.array(images_c)

        mean_centered = mu - global_mean

        S_b += (mean_centered.T @ mean_centered) * N_c

        for x in images_c:
            x_diff = (x - mu)
            S_w += x_diff.T @ x_diff

    # Added small noise for inverse
    S_w = S_w + np.eye(len(mu)) * 1e-5
```

```
    S_w_inv = np.linalg.inv(S_w)

    eigenvalues, eigenvectors = np.linalg.eigh(S_w_inv @ S_b)

    eigenvalues = np.real(eigenvalues)
    eigenvectors = np.real(eigenvectors)
    sorted_indices = np.argsort(eigenvalues)[::-1]
    eigenvalues = eigenvalues[sorted_indices]
    eigenvectors = eigenvectors[:, sorted_indices]

    return eigenvectors[:, :2]
```

4 Results and Performance Evaluation

4.1 Classification Accuracy Analysis

The table below summarizes the train and test accuracies for LDA and QDA classifiers applied to different lower-dimensional projections.

Pipeline Strategy	**LDA Train Acc**	**LDA Test Acc**	**QDA Train Acc**	**QDA Test Acc**
FDA	0.9600	0.8900	0.9667	0.9433
PCA (75% Variance)	0.9700	0.9267	0.9833	0.9467
PCA (90% Variance)	0.9900	0.9300	1.0000	0.9733
PCA (Top 2 Components)	0.9133	0.8767	0.9200	0.9200

Table 1: Performance comparison of FDA and PCA pipelines.

Analysis: As observed in the results, increasing the variance in PCA from 75% to 90% generally improves the accuracy by retaining more features, but it also increases the risk of capturing noise. FDA consistently outperforms PCA when restricted to a low number of components because it actively optimizes the axes to separate the digit classes, rather than merely maximizing data spread.

4.2 Visualizing the Feature Space

To visually compare how effectively FDA and PCA separate the digit classes, we projected the data into 2D space.

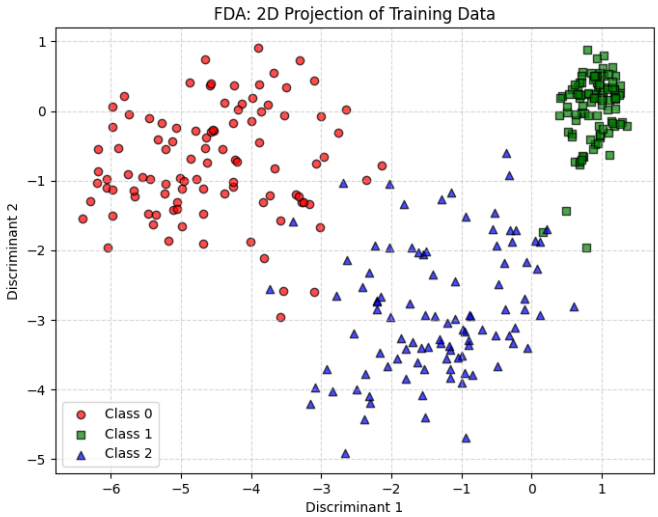

Figure 1: FDA: 2D Projection of Training Data.

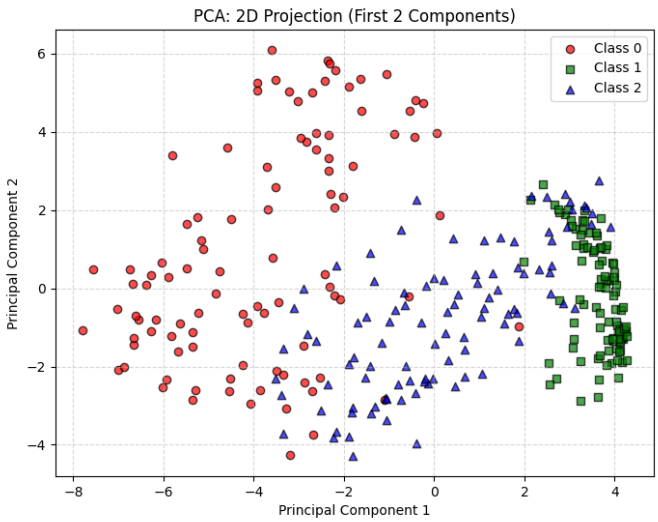

Figure 2: PCA: 2D Projection (First 2 Principal Components).

As seen in the plots above, the FDA projection results in tighter, more distinct clusters for classes 0, 1, and 2, which validates its higher classification accuracy compared to the unsupervised PCA projection.

4.3 PCA Reconstruction

To verify the integrity of the PCA transformation, we reconstructed the test samples from the 75% variance projection space back to the original 784-dimensional space by applying the inverse transformation ($YU_p^T + \mu$).

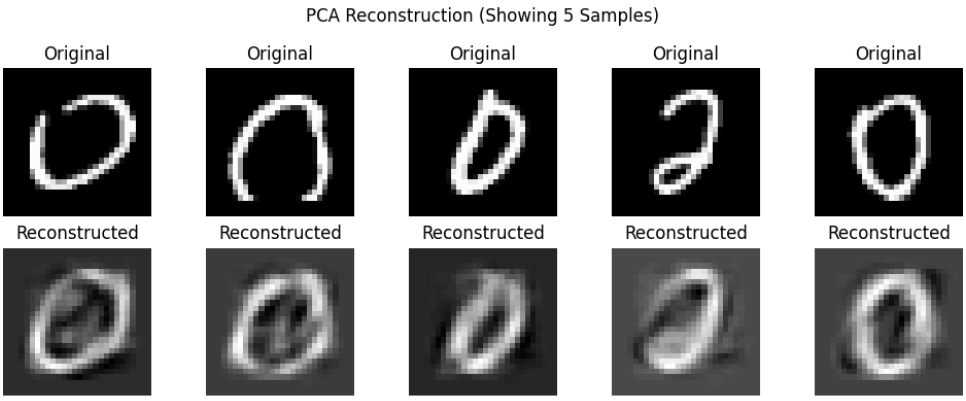

Figure 3: Original vs. Reconstructed MNIST samples (75% Variance).

The Mean Squared Error (MSE) between the original test set and the reconstructed test set was calculated to be **0.0215** While some finer details are lost, the core shapes of the digits remain easily identifiable.

5 Conclusion

This assignment successfully demonstrates the critical role of dimensionality reduction in image classification. While PCA provides an excellent mechanism for compressing data and removing noise without losing major structural variance, FDA proved to be the superior choice as a precursor to Discriminant Analysis due to its supervised nature, projecting the MNIST data specifically to maximize the boundaries between the digits 0, 1, and 2.